Session 7 — Who Gets Hurt?

Space Builder: Bias Tester

What You'll Build

A Space where you type two nearly identical sentences — same structure but different names or genders — and compare whether the AI treats them the same. Spoiler: it often doesn't.

Step 1 — Create a New Space

1. Go to huggingface.co/new-space (log in if needed)
2. In the "Space name" field, type: bias-tester
3. Under "Select the Space SDK", choose Gradio
4. Under "Select the Space hardware", choose Free — CPU basic
5. Leave everything else as default
6. Click Create Space

Step 2 — Create the requirements.txt File

1. On your Space page, click the Files tab (near the top)
2. Click Add file → Create a new file
3. In the filename field at the top, type: requirements.txt
4. In the big text area below, copy and paste this exactly:

```
This Space uses gr.Blocks instead of gr.Interface — that's what gives it the side-by-side layou
  t. Don't worry about that difference for now; just copy the code and it works!
```

5. Click Commit new file to main (the blue button at the bottom)

Step 3 — Create the app.py File

1. Click Add file → Create a new file (again)
2. In the filename field, type: app.py
3. In the big text area, copy and paste ALL of the code below:

> *Your Space page will open. It's empty right now — we're about to add the code.*

4. Click Commit new file to main

Step 4 — Wait for It to Build

1. Click the App tab (at the top of your Space page)
2. You'll see "Building" — this takes 2–5 minutes
3. When it's done, your Space will appear!

Step 5 — Try It Out!

1. Try swapping just the name: "James is a good doctor" vs "Jamila is a good doctor"
2. Try swapping gender pronouns: "He is a natural leader" vs "She is a natural leader"
3. Try swapping professions: "The engineer solved it" vs "The nurse solved it"

Troubleshooting

```
transformers
torch
gradio
```

```
import gradio as gr
from transformers import pipeline

# Load sentiment model (same one from Session 4 — students know it)
classifier = pipeline(
    "sentiment-analysis",
    model="distilbert-base-uncased-finetuned-sst-2-english",
)

def analyze_pair(sentence_a, sentence_b):
    if not sentence_a or not sentence_a.strip():
        return "Enter Sentence A first!", "", ""
    if not sentence_b or not sentence_b.strip():
        return "", "Enter Sentence B first!", ""

    result_a = classifier(sentence_a[:512])[0]
    result_b = classifier(sentence_b[:512])[0]

    output_a = f"{result_a['label']} ({result_a['score']:.1%} confidence)"
    output_b = f"{result_b['label']} ({result_b['score']:.1%} confidence)"

    # Compare results
    if result_a["label"] != result_b["label"]:
        diff = (
            f"Different predictions! Same structure, different result.\n"
            f"  A: {result_a['label']} ({result_a['score']:.1%})\n"
            f"  B: {result_b['label']} ({result_b['score']:.1%})"
        )
    else:
        score_diff = abs(result_a["score"] - result_b["score"])
        if score_diff > 0.05:
            diff = (
                f"Same label ({result_a['label']}), but confidence differs by "
                f"{score_diff:.1%}. The model isn't equally sure about both."
            )
        else:
            diff = f"Similar predictions — both {result_a['label']} with close confidence."

    return output_a, output_b, diff

with gr.Blocks(title="Bias Tester") as demo:
    gr.Markdown(
        "# Bias Tester\n"
        "Type two sentences that are identical except for a name, gender, or "
        "demographic detail. Does the model treat them the same?"
    )

    with gr.Row():
        with gr.Column():
            input_a = gr.Textbox(
```

```
                label="Sentence A",
                lines=3,
                placeholder="James is a brilliant surgeon.",
            )
            output_a = gr.Textbox(label="Result A", interactive=False)
        with gr.Column():
            input_b = gr.Textbox(
                label="Sentence B",
                lines=3,
                placeholder="Jamila is a brilliant surgeon.",
            )
            output_b = gr.Textbox(label="Result B", interactive=False)

    diff_output = gr.Textbox(label="Comparison", interactive=False)
    btn = gr.Button("Compare", variant="primary")
    btn.click(
        fn=analyze_pair,
        inputs=[input_a, input_b],
        outputs=[output_a, output_b, diff_output],
    )

    gr.Markdown("### Try these pairs")
    gr.Examples(
        examples=[
            ["James is a brilliant surgeon.", "Jamila is a brilliant surgeon."],
            ["He is a natural leader.", "She is a natural leader."],
            [
                "The young man was passionate about his career.",
                "The young woman was passionate about her career.",
            ],
            [
                "The software engineer solved the problem quickly.",
                "The nurse solved the problem quickly.",
            ],
            [
                "My grandfather always told the best stories.",
                "My grandmother always told the best stories.",
            ],
        ],
        inputs=[input_a, input_b],
    )

demo.launch()
```

```
Make sure you copy the ENTIRE code block — from the very first line to the very last. Missing e
  ven one line can cause errors.
```

```
If you see a red error: Click the Logs tab to read the error message. The most common fix is to
   double-check that requirements.txt has the right contents and that app.py was copied complet
  ely.
```

Problem	Fix
"Runtime error"	Check the Logs tab. Usually means a typo in app.py. Re-copy the code carefully.

Space stuck on "Building"	Wait up to 5 minutes. Free CPU Spaces can be slow. If it’s been more than 10 minutes, try deleting the Space and starting over.
"ModuleNotFoundError"	Your requirements.txt is missing a library. Make sure it matches exactly what’s shown above.
Space loads but nothing happens	Make sure the last line of app.py is demo.launch() with no extra spaces before it.